

Theorem 1. *With ℓ and s in $O(\text{polylog}(|V|))$, **csav** algorithm computes a k -node placement $M \subseteq V$ in time $\tilde{O}(k^2|E|)$.*

Proof. With polylogarithmically many pivots, the complexity of **rand** algorithm is $\tilde{O}(|E|)$; this is the time used to approximate penalty and generate the initial candidate S_1 . The **score** procedure takes $O(|E| + |V| \log |V|)$ and since the ℓ -confined **celf** repeats for k iterations each running **score** once, the total time of performing candidate selection in **celf-r** is $O(k(|E| + |V| \log |V|))$. Due to the limited search space, the **add-candidate** procedure evaluates marginal reward at most $k^2\ell$ times, each of which takes time $O(|E|)$. Thus we spend $O(k^2\ell \cdot |E|)$ time on updating rewards. In conclusion, the total complexity of the first phase is $\tilde{O}(k^2\ell|E| + k|V| \log |V| + |E|) = \tilde{O}(k^2|E|)$.

In the iterative improvement algorithm, each round takes $O(k(|E| + |V| \log |V|))$ to generate candidates and $O(k\ell \cdot |E|)$ to evaluate marginal reward. Thus in s rounds, it takes $O(sk(|V| \log |V| + \ell|E|)) = \tilde{O}(k|E|)$. Therefore the overall time of the algorithm is $\tilde{O}(k^2|E|)$.